

# Mastering the game of Go with deep neural networks and tree search

David Silver<sup>1\*</sup>, Aja Huang<sup>1\*</sup>, Chris J. Maddison<sup>1</sup>, Arthur Guez<sup>1</sup>, Laurent Sifre<sup>1</sup>, George van den Driessche<sup>1</sup>, Julian Schrittwieser<sup>1</sup>, Ioannis Antonoglou<sup>1</sup>, Veda Panneershelvam<sup>1</sup>, Marc Lanctot<sup>1</sup>, Sander Dieleman<sup>1</sup>, Dominik Grewe<sup>1</sup>, John Nham<sup>2</sup>, Nal Kalchbrenner<sup>1</sup>, Ilya Sutskever<sup>2</sup>, Timothy Lillicrap<sup>1</sup>, Madeleine Leach<sup>1</sup>, Koray Kavukcuoglu<sup>1</sup>, Thore Graepel<sup>1</sup> & Demis Hassabis<sup>1</sup>

The game of Go has long been viewed as the most challenging of classic games for artificial intelligence owing to its enormous search space and the difficulty of evaluating board positions and moves. Here we introduce a new approach to computer Go that uses ‘value networks’ to evaluate board positions and ‘policy networks’ to select moves. These deep neural networks are trained by a novel combination of supervised learning from human expert games, and reinforcement learning from games of self-play. Without any lookahead search, the neural networks play Go at the level of state-of-the-art Monte Carlo tree search programs that simulate thousands of random games of self-play. We also introduce a new search algorithm that combines Monte Carlo simulation with value and policy networks. Using this search algorithm, our program AlphaGo achieved a 99.8% winning rate against other Go programs, and defeated the human European Go champion by 5 games to 0. This is the first time that a computer program has defeated a human professional player in the full-sized game of Go, a feat previously thought to be at least a decade away.

All games of perfect information have an optimal value function,  $v^*(s)$ , which determines the outcome of the game, from every board position or state  $s$ , under perfect play by all players. These games may be solved by recursively computing the optimal value function in a search tree containing approximately  $b^d$  possible sequences of moves, where  $b$  is the game’s breadth (number of legal moves per position) and  $d$  is its depth (game length). In large games, such as chess ( $b \approx 35$ ,  $d \approx 80$ )<sup>1</sup> and especially Go ( $b \approx 250$ ,  $d \approx 150$ )<sup>1</sup>, exhaustive search is infeasible<sup>2,3</sup>, but the effective search space can be reduced by two general principles. First, the depth of the search may be reduced by position evaluation: truncating the search tree at state  $s$  and replacing the subtree below  $s$  by an approximate value function  $v(s) \approx v^*(s)$  that predicts the outcome from state  $s$ . This approach has led to superhuman performance in chess<sup>4</sup>, checkers<sup>5</sup> and othello<sup>6</sup>, but it was believed to be intractable in Go due to the complexity of the game<sup>7</sup>. Second, the breadth of the search may be reduced by sampling actions from a policy  $p(a|s)$  that is a probability distribution over possible moves  $a$  in position  $s$ . For example, Monte Carlo rollouts<sup>8</sup> search to maximum depth without branching at all, by sampling long sequences of actions for both players from a policy  $p$ . Averaging over such rollouts can provide an effective position evaluation, achieving superhuman performance in backgammon<sup>8</sup> and Scrabble<sup>9</sup>, and weak amateur level play in Go<sup>10</sup>.

Monte Carlo tree search (MCTS)<sup>11,12</sup> uses Monte Carlo rollouts to estimate the value of each state in a search tree. As more simulations are executed, the search tree grows larger and the relevant values become more accurate. The policy used to select actions during search is also improved over time, by selecting children with higher values. Asymptotically, this policy converges to optimal play, and the evaluations converge to the optimal value function<sup>12</sup>. The strongest current Go programs are based on MCTS, enhanced by policies that are trained to predict human expert moves<sup>13</sup>. These policies are used to narrow the search to a beam of high-probability actions, and to sample actions during rollouts. This approach has achieved strong amateur play<sup>13–15</sup>. However, prior work has been limited to shallow

然而，之前的工作仅限于浅层策略或基于输入特征的线性组合的价值函数。

policies<sup>13–15</sup> or value functions<sup>16</sup> based on a linear combination of input features.

Recently, deep convolutional neural networks have achieved unprecedented performance in visual domains: for example, image classification<sup>17</sup>, face recognition<sup>18</sup>, and playing Atari games<sup>19</sup>. They use many layers of neurons, each arranged in overlapping tiles, to construct increasingly abstract, localized representations of an image<sup>20</sup>. We employ a similar architecture for the game of Go. We pass in the board position as a  $19 \times 19$  image and use convolutional layers to construct a representation of the position. We use these neural networks to reduce the effective depth and breadth of the search tree: evaluating positions using a value network, and sampling actions using a policy network.

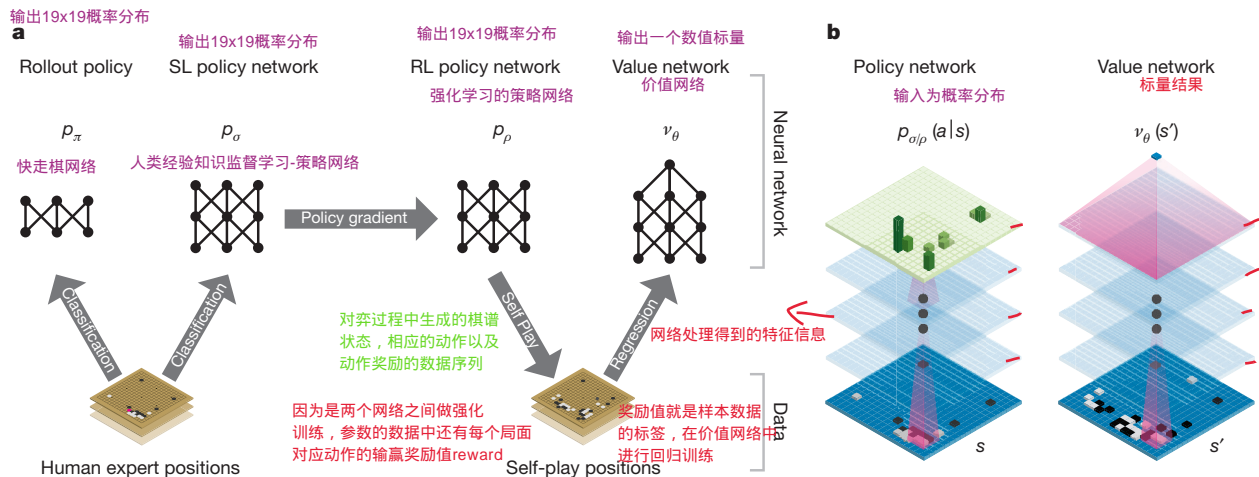
We train the neural networks using a pipeline consisting of several stages of machine learning (Fig. 1). We begin by training a supervised learning (SL) policy network  $p_\pi$  directly from expert human moves. This provides fast, efficient learning updates with immediate feedback and high-quality gradients. Similar to prior work<sup>13,15</sup>, we also train a fast policy  $p_\pi$  that can rapidly sample actions during rollouts. Next, we train a reinforcement learning (RL) policy network  $p_\theta$  that improves the SL policy network by optimizing the final outcome of games of self-play. This adjusts the policy towards the correct goal of winning games, rather than maximizing predictive accuracy. Finally, we train a value network  $v_\theta$  that predicts the winner of games played by the RL policy network against itself. Our program AlphaGo efficiently combines the policy and value networks with MCTS.

## Supervised learning of policy networks

For the first stage of the training pipeline, we build on prior work on predicting expert moves in the game of Go using supervised learning<sup>13,21–24</sup>. The SL policy network  $p_\pi(a|s)$  alternates between convolutional layers with weights  $\sigma$ , and rectifier nonlinearities. A final softmax layer outputs a probability distribution over all legal moves  $a$ . The input  $s$  to the policy network is a simple representation of the board state (see Extended Data Table 2). The policy network is trained on randomly

<sup>1</sup>Google DeepMind, 5 New Street Square, London EC4A 3TW, UK. <sup>2</sup>Google, 1600 Amphitheatre Parkway, Mountain View, California 94043, USA.

\*These authors contributed equally to this work.



**Figure 1 | Neural network training pipeline and architecture.** **a**, A fast rollout policy  $p_\pi$  and supervised learning (SL) policy network  $p_\sigma$  are trained to predict human expert moves in a data set of positions. A reinforcement learning (RL) policy network  $p_\rho$  is initialized to the SL policy network, and is then improved by policy gradient learning to maximize the outcome (that is, winning more games) against previous versions of the policy network. A new data set is generated by playing games of self-play with the RL policy network. Finally, a value network  $v_\theta$  is trained by regression to predict the expected outcome (that is, whether

通过回归训练一个价值网络  $v$ ，以预测来自自我游戏数据集的位置的预期结果。 sampled state-action pairs  $(s, a)$ , using stochastic gradient ascent to maximize the likelihood of the human move  $a$  selected in state  $s$


$$\Delta\sigma \propto \frac{\partial \log p_\sigma(a|s)}{\partial \sigma}$$

，数学符号，表示与某个量成正比。可以表示为  $y \propto x$

We trained a 13-layer policy network, which we call the SL policy network, from 30 million positions from the KGS Go Server. The network predicted expert moves on a held out test set with an accuracy of 57.0% using all input features, and 55.7% using only raw board position and move history as inputs, compared to the state-of-the-art from other research groups of 44.4% at date of submission<sup>24</sup> (full results in Extended Data Table 3). Small improvements in accuracy led to large improvements in playing strength (Fig. 2a); larger networks achieve better accuracy but are slower to evaluate during search. We also trained a faster but less accurate rollout policy  $p_\pi(a|s)$ , using a linear softmax of small pattern features (see Extended Data Table 4) with weights  $\pi$ ; this achieved an accuracy of 24.2%, using just 2  $\mu$ s to select an action, rather than 3 ms for the policy network.

## Reinforcement learning of policy networks

The second stage of the training pipeline aims at improving the policy network by policy gradient reinforcement learning (RL)<sup>25,26</sup>. The RL policy network  $p_\rho$  is identical in structure to the SL policy network,



**Figure 2 | Strength and accuracy of policy and value networks.** **a**, Plot showing the playing strength of policy networks as a function of their training accuracy. Policy networks with 128, 192, 256 and 384 convolutional filters per layer were evaluated periodically during training; the plot shows the winning rate of AlphaGo using that policy network against the match version of AlphaGo. **b**, Comparison of evaluation accuracy between the value network and rollouts with different policies.


the current player wins) in positions from the self-play data set.

**b**, Schematic representation of the neural network architecture used in AlphaGo. The policy network takes a representation of the board position  $s$  as its input, passes it through many convolutional layers with parameters  $\sigma$  (SL policy network) or  $\rho$  (RL policy network), and outputs a probability distribution  $p_\sigma(a|s)$  or  $p_\rho(a|s)$  over legal moves  $a$ , represented by a probability map over the board. The value network similarly uses many convolutional layers with parameters  $\theta$ , but outputs a scalar value  $v_\theta(s')$  that predicts the expected outcome in position  $s'$ .

and its weights  $\rho$  are initialized to the same values,  $\rho = \sigma$ . (We play games between the current policy network  $p_\rho$  and a randomly selected previous iteration of the policy network.) Randomizing from a pool of opponents in this way stabilizes training by preventing overfitting to the current policy. We use a reward function  $r(s)$  that is zero for all non-terminal time steps  $t < T$ . The outcome  $z_t = \pm 1$  is the terminal reward at the end of the game from the perspective of the current player at time step  $t$ : +1 for winning and -1 for losing. Weights are then updated at each time step  $t$  by stochastic gradient ascent in the direction that maximizes expected outcome<sup>25</sup>

$$\Delta\rho \propto \frac{\partial \log p_\rho(a_t|s_t)}{\partial \rho} z_t$$

We evaluated the performance of the RL policy network in game play, sampling each move  $a_t \sim p_\rho(\cdot|s_t)$  from its output probability distribution over actions. When played head-to-head, the RL policy network won more than 80% of games against the SL policy network. We also tested against the strongest open-source Go program, Pachi<sup>14</sup>, a sophisticated Monte Carlo search program, ranked at 2 amateur *dan* on KGS, that executes 100,000 simulations per move. Using no search at all, the RL policy network won 85% of games against Pachi. In comparison, the previous state-of-the-art, based only on supervised




positions and outcomes were sampled from human expert games. Each position was evaluated by a single forward pass of the value network  $v_\theta$ , or by the mean outcome of 100 rollouts, played out using either uniform random rollouts, the fast rollout policy  $p_\pi$ , the SL policy network  $p_\sigma$  or the RL policy network  $p_\rho$ . The mean squared error between the predicted value and the actual game outcome is plotted against the stage of the game (how many moves had been played in the given position).

这一步，用策略网络提前计算出该状态s下的动作概率分布，就是给这个状态s下计算好未来可能动作的先验概率  
当未来的某一步真的以该状态节点做根节点，就可以直接使用先验概率了，去完成MCTS迭代

把模拟进行的每一个随机棋局称为一次推演(rollout)

这里都是给了某个状态s了，当是从空状态开始呢？其实就是把当前空状态输入给网络，得到候选边，去完成MCTS迭代



叶子节点可以被扩展；新的节点被策略网络 $p_\theta$ 处理一次。由策略网络 $p_\theta$ 处理，输出概率被存储为每个行动的先验概率 $P$ 。在模拟结束时，叶子节点以两种方式进行评估：使用价值网络 $v_\theta$ ；以及通过快速滚动策略 $p_\pi$ 运行到游戏结束，然后用函数 $r$ 计算出赢家。行动值 $Q$ 被更新以跟踪该行动下方子树中所有评估的平均值 $r(\cdot)$ 和 $v_\theta(\cdot)$ 。

**Figure 3 | Monte Carlo tree search in AlphaGo.** a, Each simulation traverses the tree by selecting the edge with maximum action value  $Q$ , plus a bonus  $u(P)$  that depends on a stored prior probability  $P$  for that edge. b, The leaf node may be expanded; the new node is processed once by the policy network  $p_\theta$ , and the output probabilities are stored as prior probabilities  $P$  for each action. c, At the end of a simulation, the leaf node

learning of convolutional networks, won 11% of games against Pachi<sup>23</sup> and 12% against a slightly weaker program, Fuego<sup>24</sup>.

### Reinforcement learning of value networks

The final stage of the training pipeline focuses on position evaluation, estimating a value function  $v^\theta(s)$  that predicts the outcome from position  $s$  of games played by using policy  $p$  for both players<sup>28–30</sup>

$$v^\theta(s) = \mathbb{E}[z_t | s_t = s, a_{t:T} \sim p]$$

Ideally, we would like to know the optimal value function under perfect play  $v^*(s)$ ; in practice, we instead estimate the value function  $v_{p_\theta}$  for our strongest policy, using the RL policy network  $p_\theta$ . We approximate the value function using a value network  $v_\theta(s)$  with weights  $\theta$ ,  $v_\theta(s) \approx v_{p_\theta}(s) \approx v^*(s)$ . This neural network has a similar architecture to the policy network, but outputs a single prediction instead of a probability distribution. We train the weights of the value network by regression on state-outcome pairs  $(s, z)$ , using stochastic gradient descent to minimize the mean squared error (MSE) between the predicted value  $v_\theta(s)$ , and the corresponding outcome  $z$ .

在进行强化训练策略网络时，两个博弈就会产生训练数据包好状态-动作-奖励值 $(s, a, r, t)$  (s,z)就是 $(s, t, r, t)$

$$\Delta \theta \propto \frac{\partial v_\theta(s)}{\partial \theta} (z - v_\theta(s))$$

The naive approach of predicting game outcomes from data consisting of complete games leads to overfitting. The problem is that successive positions are strongly correlated, differing by just one stone, but the regression target is shared for the entire game. When trained on the KGS data set in this way, the value network memorized the game outcomes rather than generalizing to new positions, achieving a minimum MSE of 0.37 on the test set, compared to 0.19 on the training set. To mitigate this problem, we generated a new self-play data set consisting of 30 million distinct positions, each sampled from a separate game. Each game was played between the RL policy network and itself until the game terminated. Training on this data set led to MSEs of 0.226 and 0.234 on the training and test set respectively, indicating minimal overfitting. Figure 2b shows the position evaluation accuracy of the value network, compared to Monte Carlo rollouts using the fast rollout policy  $p_\pi$ ; the value function was consistently more accurate. A single evaluation of  $v_\theta(s)$  also approached the accuracy of Monte Carlo rollouts using the RL policy network  $p_\theta$ , but using 15,000 times less computation.

### Searching with policy and value networks

AlphaGo combines the policy and value networks in an MCTS algorithm (Fig. 3) that selects actions by lookahead search. Each edge AlphaGo用MCTS来把策略网络和价值网络进行组织在一起

is evaluated in two ways: using the value network  $v_\theta$ ; and by running a rollout to the end of the game with the fast rollout policy  $p_\pi$ , then computing the winner with function  $r$ . d, Action values  $Q$  are updated to track the mean value of all evaluations  $r(\cdot)$  and  $v_\theta(\cdot)$  in the subtree below that action.

Q(): 对应状态s,采取动作a的价值

(state,action) (局面状态,采用动作)

$(s, a)$  of the search tree stores an action value  $Q(s, a)$ , visit count  $N(s, a)$ , and prior probability  $P(s, a)$ . The tree is traversed by simulation (that is, descending the tree in complete games without backup), starting from the root state. At each time step  $t$  of each simulation, an action  $a_t$  is selected from state  $s_t$ . 采取的动作 $a_t$ 是具有最大价值的和最高访问的节点

$$a_t = \operatorname{argmax}_a (Q(s_t, a) + u(s_t, a))$$

最大行动价值，再加上一个奖励

so as to maximize action value plus a bonus

遍历节点的函数，正比于先验概率，反比访问的次数

$$u(s, a) \propto \frac{P(s, a)}{1 + N(s, a)}$$

与先验概率成正比，但随着重复访问而衰减，以鼓励探索。

that is proportional to the prior probability but decays with repeated visits to encourage exploration. When the traversal reaches a leaf node  $s_L$  at step  $L$ , the leaf node may be expanded. The leaf position  $s_L$  is processed just once by the SL policy network  $p_\theta$ . The output probabilities are stored as prior probabilities  $P$  for each legal action  $a$ ,  $P(s, a) = p_\theta(a | s)$ . The leaf node is evaluated in two very different ways: first, by the value network  $v_\theta(s_L)$ ; and second, by the outcome  $z_L$  of a random rollout played out until terminal step  $T$  using the fast rollout policy  $p_\pi$ ; these evaluations are combined, using a mixing parameter  $\lambda$ , into a leaf evaluation  $V(s_L)$ .

$$v_i(s) = \lambda \text{value}(s) + (1 - \lambda) \text{rollout}(s) \\ V(s_L) = (1 - \lambda) v_\theta(s_L) + \lambda z_L$$

At the end of simulation, the action values and visit counts of all traversed edges are updated. Each edge accumulates the visit count and mean evaluation of all simulations passing through that edge. 更新树的每条边的数据情况

$$N(s, a) = \sum_{i=1}^n 1(s, a, i)$$

$$Q(s, a) = \frac{1}{N(s, a)} \sum_{i=1}^n 1(s, a, i) V(s_L^i)$$


select leaf i 第i次模拟的叶子节点  
(s,a,i):应该是个0/1值，表示是否遍历了这个树的边

where  $s_L^i$  is the leaf node from the  $i$ th simulation, and  $1(s, a, i)$  indicates whether an edge  $(s, a)$  was traversed during the  $i$ th simulation. Once the search is complete, the algorithm chooses the most visited move from the root position.

It is worth noting that the SL policy network  $p_\theta$  performed better in AlphaGo than the stronger RL policy network  $p_\pi$ , presumably because humans select a diverse beam of promising moves, whereas RL optimizes for the single best move. However, the value function  $v_\theta(s) \approx v_{p_\theta}(s)$  derived from the stronger RL policy network performed

其中， $s_L$ 是第i次模拟的叶子节点， $1(s, a, i)$ 表示在第i次模拟中是否穿越了一条边 $(s, a)$ 。一旦搜索完成，该算法就会选择从根部位置开始的最多访问的移动。





**Figure 4 | Tournament evaluation of AlphaGo.** **a**, Results of a tournament between different Go programs (see Extended Data Tables 6–11). Each program used approximately 5 s computation time per move. To provide a greater challenge to AlphaGo, some programs (pale upper bars) were given four handicap stones (that is, free moves at the start of every game) against all opponents. Programs were evaluated on an Elo scale<sup>37</sup>: a 230 point gap corresponds to a 79% probability of winning, which roughly corresponds to one amateur *dan* rank advantage on KGS<sup>38</sup>; an approximate correspondence to human ranks is also shown,

horizontal lines show KGS ranks achieved online by that program. Games against the human European champion Fan Hui were also included; these games used longer time controls. 95% confidence intervals are shown. **b**, Performance of AlphaGo, on a single machine, for different combinations of components. The version solely using the policy network does not perform any search. **c**, Scalability study of MCTS in AlphaGo with search threads and GPUs, using asynchronous search (light blue) or distributed search (dark blue), for 2 s per move.

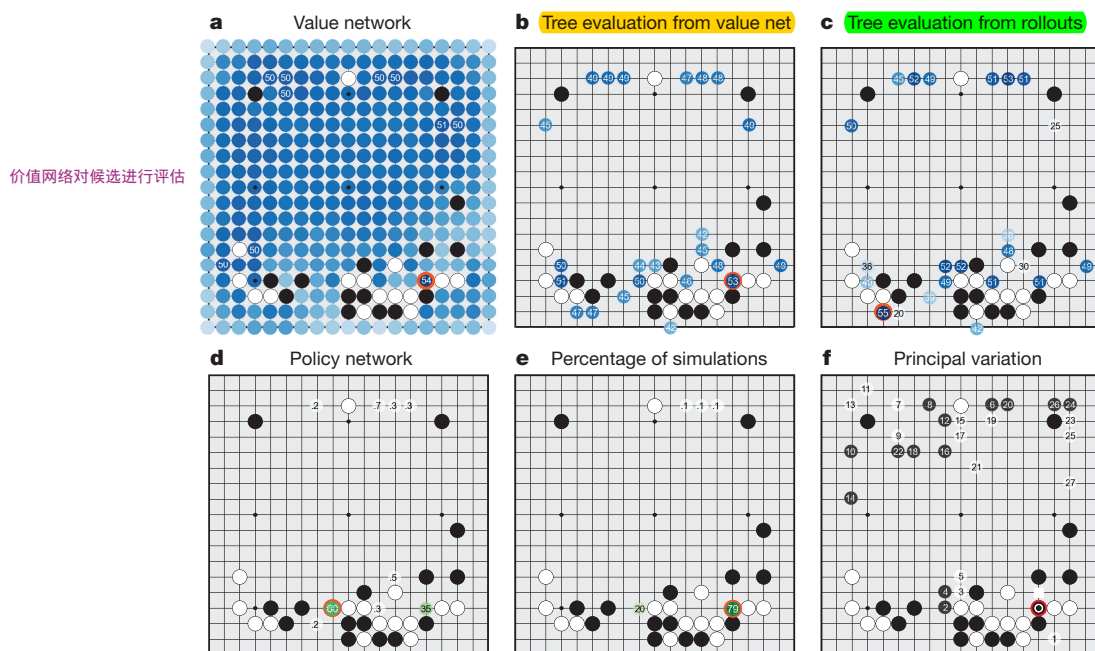
better in AlphaGo than a value function  $v_\theta(s) \approx v^{p_\theta}(s)$  derived from the SL policy network.

Evaluating policy and value networks requires several orders of magnitude more computation than traditional search heuristics. To efficiently combine MCTS with deep neural networks, AlphaGo uses an asynchronous multi-threaded search that executes simulations on CPUs, and computes policy and value networks in parallel on GPUs. The final version of AlphaGo used 40 search threads, 48 CPUs, and 8 GPUs. We also implemented a distributed version of AlphaGo that

exploited multiple machines, 40 search threads, 1,202 CPUs and 176 GPUs. The Methods section provides full details of asynchronous and distributed MCTS.


### Evaluating the playing strength of AlphaGo

To evaluate AlphaGo, we ran an internal tournament among variants of AlphaGo and several other Go programs, including the strongest commercial programs Crazy Stone<sup>13</sup> and Zen, and the strongest open source programs Pachi<sup>14</sup> and Fuego<sup>15</sup>. All of these programs are based



**Figure 5 | How AlphaGo (black, to play) selected its move in an informal game against Fan Hui.** For each of the following statistics, the location of the maximum value is indicated by an orange circle. **a**, Evaluation of all successors  $s'$  of the root position  $s$ , using the value network  $v_\theta(s')$ ; estimated winning percentages are shown for the top evaluations. **b**, Action values  $Q(s, a)$  for each edge  $(s, a)$  in the tree from root position  $s$ ; averaged over value network evaluations only ( $\lambda = 0$ ). **c**, Action values  $Q(s, a)$ , averaged over rollout evaluations only ( $\lambda = 1$ ).

**d**, Move probabilities directly from the SL policy network,  $p_\theta(a|s)$ ; reported as a percentage (if above 0.1%). **e**, Percentage frequency with which actions were selected from the root during simulations. **f**, The principal variation (path with maximum visit count) from AlphaGo's search tree. The moves are presented in a numbered sequence. AlphaGo selected the move indicated by the red circle; Fan Hui responded with the move indicated by the white square; in his post-game commentary he preferred the move (labelled 1) predicted by AlphaGo.



**Figure 6 | Games from the match between AlphaGo and the European champion, Fan Hui.** Moves are shown in a numbered sequence corresponding to the order in which they were played. Repeated moves on the same intersection are shown in pairs below the board. The first

move number in each pair indicates when the repeat move was played, at an intersection identified by the second move number (see Supplementary Information).

on high-performance MCTS algorithms. In addition, we included the open source program GnuGo, a Go program using state-of-the-art search methods that preceded MCTS. All programs were allowed 5 s of computation time per move.

The results of the tournament (see Fig. 4a) suggest that single-machine AlphaGo is many *dan* ranks stronger than any previous Go program, winning 494 out of 495 games (99.8%) against other Go programs. To provide a greater challenge to AlphaGo, we also played games with four handicap stones (that is, free moves for the opponent); AlphaGo won 77%, 86%, and 99% of handicap games against Crazy Stone, Zen and Pachi, respectively. The distributed version of AlphaGo was significantly stronger, winning 77% of games against single-machine AlphaGo and 100% of its games against other programs.

We also assessed variants of AlphaGo that evaluated positions using just the value network ( $\lambda = 0$ ) or just rollouts ( $\lambda = 1$ ) (see Fig. 4b). Even without rollouts AlphaGo exceeded the performance of all other Go programs, demonstrating that value networks provide a viable alternative to Monte Carlo evaluation in Go. However, the mixed evaluation ( $\lambda = 0.5$ ) performed best, winning  $\geq 95\%$  of games against other variants. This suggests that the two position-evaluation

mechanisms are complementary: the value network approximates the outcome of games played by the strong but impractically slow  $p_\rho$ , while the rollouts can precisely score and evaluate the outcome of games played by the weaker but faster rollout policy  $p_\pi$ . Figure 5 visualizes the evaluation of a real game position by AlphaGo.

Finally, we evaluated the distributed version of AlphaGo against Fan Hui, a professional 2 *dan*, and the winner of the 2013, 2014 and 2015 European Go championships. Over 5–9 October 2015 AlphaGo and Fan Hui competed in a formal five-game match. AlphaGo won the match 5 games to 0 (Fig. 6 and Extended Data Table 1). This is the first time that a computer Go program has defeated a human professional player, without handicap, in the full game of Go—a feat that was previously believed to be at least a decade away<sup>3,7,31</sup>.

## Discussion

In this work we have developed a Go program, based on a combination of deep neural networks and tree search, that plays at the level of the strongest human players, thereby achieving one of artificial intelligence's “grand challenges”<sup>31–33</sup>. We have developed, for the first time, effective move selection and position evaluation functions for Go, based on deep neural networks that are trained by a novel combination

of supervised and reinforcement learning. We have introduced a new search algorithm that successfully combines neural network evaluations with Monte Carlo rollouts. Our program AlphaGo integrates these components together, at scale, in a high-performance tree search engine.

During the match against Fan Hui, AlphaGo evaluated thousands of times fewer positions than Deep Blue did in its chess match against Kasparov<sup>4</sup>; compensating by selecting those positions more intelligently, using the policy network, and evaluating them more precisely, using the value network—an approach that is perhaps closer to how humans play. Furthermore, while Deep Blue relied on a handcrafted evaluation function, the neural networks of AlphaGo are trained directly from gameplay purely through general-purpose supervised and reinforcement learning methods.

Go is exemplary in many ways of the difficulties faced by artificial intelligence<sup>33,34</sup>: a challenging decision-making task, an intractable search space, and an optimal solution so complex it appears infeasible to directly approximate using a policy or value function. The previous major breakthrough in computer Go, the introduction of MCTS, led to corresponding advances in many other domains; for example, general game-playing, classical planning, partially observed planning, scheduling, and constraint satisfaction<sup>35,36</sup>. By combining tree search with policy and value networks, AlphaGo has finally reached a professional level in Go, providing hope that human-level performance can now be achieved in other seemingly intractable artificial intelligence domains.

**Online Content** Methods, along with any additional Extended Data display items and Source Data, are available in the online version of the paper; references unique to these sections appear only in the online paper.

**Received 11 November 2015; accepted 5 January 2016.**

- Allis, L. V. *Searching for Solutions in Games and Artificial Intelligence*. PhD thesis, Univ. Limburg, Maastricht, The Netherlands (1994).
- van den Herik, H., Uiterwijk, J. W. & van Rijswijk, J. Games solved: now and in the future. *Artif. Intell.* **134**, 277–311 (2002).
- Schaeffer, J. The games computers (and people) play. *Advances in Computers* **52**, 189–266 (2000).
- Campbell, M., Hoane, A. & Hsu, F. Deep Blue. *Artif. Intell.* **134**, 57–83 (2002).
- Schaeffer, J. et al. A world championship caliber checkers program. *Artif. Intell.* **53**, 273–289 (1992).
- Buro, M. From simple features to sophisticated evaluation functions. In *1st International Conference on Computers and Games*, 126–145 (1999).
- Müller, M. Computer Go. *Artif. Intell.* **134**, 145–179 (2002).
- Tesauro, G. & Galperin, G. On-line policy improvement using Monte-Carlo search. In *Advances in Neural Information Processing*, 1068–1074 (1996).
- Sheppard, B. World-championship-caliber Scrabble. *Artif. Intell.* **134**, 241–275 (2002).
- Bouzy, B. & Helmstetter, B. Monte-Carlo Go developments. In *10th International Conference on Advances in Computer Games*, 159–174 (2003).
- Coulom, R. Efficient selectivity and backup operators in Monte-Carlo tree search. In *5th International Conference on Computers and Games*, 72–83 (2006).
- Kocsis, L. & Szepesvári, C. Bandit based Monte-Carlo planning. In *15th European Conference on Machine Learning*, 282–293 (2006).
- Coulom, R. Computing Elo ratings of move patterns in the game of Go. *ICGA J.* **30**, 198–208 (2007).
- Baudiš, P. & Gailly, J.-L. Pachi: State of the art open source Go program. In *Advances in Computer Games*, 24–38 (Springer, 2012).
- Müller, M., Enzenberger, M., Arneson, B. & Segal, R. Fuego – an open-source framework for board games and Go engine based on Monte-Carlo tree search. *IEEE Trans. Comput. Intell. AI in Games* **2**, 259–270 (2010).
- Gelly, S. & Silver, D. Combining online and offline learning in UCT. In *17th International Conference on Machine Learning*, 273–280 (2007).
- Krizhevsky, A., Sutskever, I. & Hinton, G. ImageNet classification with deep convolutional neural networks. In *Advances in Neural Information Processing Systems*, 1097–1105 (2012).
- Lawrence, S., Giles, C. L., Tsoi, A. C. & Back, A. D. Face recognition: a convolutional neural-network approach. *IEEE Trans. Neural Netw.* **8**, 98–113 (1997).
- Mnih, V. et al. Human-level control through deep reinforcement learning. *Nature* **518**, 529–533 (2015).
- LeCun, Y., Bengio, Y. & Hinton, G. Deep learning. *Nature* **521**, 436–444 (2015).
- Stern, D., Herbrich, R. & Graepel, T. Bayesian pattern ranking for move prediction in the game of Go. In *International Conference of Machine Learning*, 873–880 (2006).
- Sutskever, I. & Nair, V. Mimicking Go experts with convolutional neural networks. In *International Conference on Artificial Neural Networks*, 101–110 (2008).
- Maddison, C. J., Huang, A., Sutskever, I. & Silver, D. Move evaluation in Go using deep convolutional neural networks. *3rd International Conference on Learning Representations* (2015).
- Clark, C. & Storkey, A. J. Training deep convolutional neural networks to play go. In *32nd International Conference on Machine Learning*, 1766–1774 (2015).
- Williams, R. J. Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Mach. Learn.* **8**, 229–256 (1992).
- Sutton, R., McAllester, D., Singh, S. & Mansour, Y. Policy gradient methods for reinforcement learning with function approximation. In *Advances in Neural Information Processing Systems*, 1057–1063 (2000).
- Sutton, R. & Barto, A. *Reinforcement Learning: an Introduction* (MIT Press, 1998).
- Schraudolph, N. N., Dayan, P. & Sejnowski, T. J. Temporal difference learning of position evaluation in the game of Go. *Adv. Neural Inf. Process. Syst.* **6**, 817–824 (1994).
- Enzenberger, M. Evaluation in Go by a neural network using soft segmentation. In *10th Advances in Computer Games Conference*, 97–108 (2003).
- Silver, D., Sutton, R. & Müller, M. Temporal-difference search in computer Go. *Mach. Learn.* **87**, 183–219 (2012).
- Levinovitz, A. The mystery of Go, the ancient game that computers still can't win. *Wired Magazine* (2014).
- Mechner, D. All Systems Go. *The Sciences* **38**, 32–37 (1998).
- Mandziuk, J. Computational intelligence in mind games. In *Challenges for Computational Intelligence*, 407–442 (2007).
- Berliner, H. A chronology of computer chess and its literature. *Artif. Intell.* **10**, 201–214 (1978).
- Browne, C. et al. A survey of Monte-Carlo tree search methods. *IEEE Trans. Comput. Intell. AI in Games* **4**, 1–43 (2012).
- Gelly, S. et al. The grand challenge of computer Go: Monte Carlo tree search and extensions. *Commun. ACM* **55**, 106–113 (2012).
- Coulom, R. Whole-history rating: A Bayesian rating system for players of time-varying strength. In *International Conference on Computers and Games*, 113–124 (2008).
- KGS. Rating system math. <http://www.gokgs.com/help/rmath.html>.

**Supplementary Information** is available in the online version of the paper.

**Acknowledgements** We thank Fan Hui for agreeing to play against AlphaGo; T. Manning for refereeing the match; R. Munos and T. Schaul for helpful discussions and advice; A. Cain and M. Cant for work on the visuals; P. Dayan, G. Wayne, D. Kumaran, D. Purves, H. van Hasselt, A. Barreto and G. Ostrovski for reviewing the paper; and the rest of the DeepMind team for their support, ideas and encouragement.

**Author Contributions** A.H., G.v.d.D., J.S., I.A., M.La., A.G., T.G. and D.S. designed and implemented the search in AlphaGo. C.J.M., A.G., L.S., A.H., I.A., V.P., S.D., D.G., N.K., I.S., K.K. and D.S. designed and trained the neural networks in AlphaGo. J.S., J.N., A.H. and D.S. designed and implemented the evaluation framework for AlphaGo. D.S., M.Le., T.L., T.G., K.K. and D.H. managed and advised on the project. D.S., T.G., A.G. and D.H. wrote the paper.

**Author Information** Reprints and permissions information is available at [www.nature.com/reprints](http://www.nature.com/reprints). The authors declare no competing financial interests. Readers are welcome to comment on the online version of the paper. Correspondence and requests for materials should be addressed to D.S. ([davidsilver@google.com](mailto:davidsilver@google.com)) or D.H. ([demishassabis@google.com](mailto:demishassabis@google.com)).



## METHODS

**Problem setting.** Many games of perfect information, such as chess, checkers, othello, backgammon and Go, may be defined as alternating Markov games<sup>39</sup>. In these games, there is a state space  $\mathcal{S}$  (where state includes an indication of the current player to play); an action space  $\mathcal{A}(s)$  defining the legal actions in any given state  $s \in \mathcal{S}$ ; a state transition function  $f(s, a, \xi)$  defining the successor state after selecting action  $a$  in state  $s$  and random input  $\xi$  (for example, dice); and finally a reward function  $r^i(s)$  describing the reward received by player  $i$  in state  $s$ . We restrict our attention to two-player zero-sum games,  $r^1(s) = -r^2(s) = r(s)$ , with deterministic state transitions,  $f(s, a, \xi) = f(s, a)$ , and zero rewards except at a terminal time step  $T$ . The outcome of the game  $z_t = \pm r(s_T)$  is the terminal reward at the end of the game from the perspective of the current player at time step  $t$ . A policy  $p(a|s)$  is a probability distribution over legal actions  $a \in \mathcal{A}(s)$ . A value function is the expected outcome if all actions for both players are selected according to policy  $p$ , that is,  $v^p(s) = \mathbb{E}[z_t | s_t = s, a_{t:T} \sim p]$ . Zero-sum games have a unique optimal value function  $v^*(s)$  that determines the outcome from state  $s$  following perfect play by both players,

$$v^*(s) = \begin{cases} z_T & \text{if } s = s_T, \\ \max_a - v^*(f(s, a)) & \text{otherwise} \end{cases}$$

**Prior work.** The optimal value function can be computed recursively by minimax (or equivalently negamax) search<sup>40</sup>. Most games are too large for exhaustive minimax tree search; instead, the game is truncated by using an approximate value function  $v(s) \approx v^*(s)$  in place of terminal rewards. Depth-first minimax search with alpha-beta pruning<sup>40</sup> has achieved superhuman performance in chess<sup>4</sup>, checkers<sup>5</sup> and othello<sup>6</sup>, but it has not been effective in Go<sup>7</sup>.

Reinforcement learning can learn to approximate the optimal value function directly from games of self-play<sup>39</sup>. The majority of prior work has focused on a linear combination  $v_\theta(s) = \varphi(s) \cdot \theta$  of features  $\varphi(s)$  with weights  $\theta$ . Weights were trained using temporal-difference learning<sup>41</sup> in chess<sup>42,43</sup>, checkers<sup>44,45</sup> and Go<sup>30</sup>; or using linear regression in othello<sup>6</sup> and Scrabble<sup>9</sup>. Temporal-difference learning has also been used to train a neural network to approximate the optimal value function, achieving superhuman performance in backgammon<sup>46</sup>; and achieving weak *kyu*-level performance in small-board Go<sup>28,29,47</sup> using convolutional networks.

An alternative approach to minimax search is Monte Carlo tree search (MCTS)<sup>11,12</sup>, which estimates the optimal value of interior nodes by a double approximation,  $V^n(s) \approx v^n(s) \approx v^*(s)$ . The first approximation,  $V^n(s) \approx v^n(s)$ , uses  $n$  Monte Carlo simulations to estimate the value function of a simulation policy  $P^n$ . The second approximation,  $v^n(s) \approx v^*(s)$ , uses a simulation policy  $P^n$  in place of minimax optimal actions. The simulation policy selects actions according to a search control function  $\arg\max_a (Q^n(s, a) + u(s, a))$ , such as UCT<sup>12</sup>, that selects children with higher action values,  $Q^n(s, a) = -V^n(f(s, a))$ , plus a bonus  $u(s, a)$  that encourages exploration; or in the absence of a search tree at state  $s$ , it samples actions from a fast rollout policy  $p_n(a|s)$ . As more simulations are executed and the search tree grows deeper, the simulation policy becomes informed by increasingly accurate statistics. In the limit, both approximations become exact and MCTS (for example, with UCT) converges<sup>12</sup> to the optimal value function  $\lim_{n \rightarrow \infty} V^n(s) = \lim_{n \rightarrow \infty} v^n(s) = v^*(s)$ . The strongest current Go programs are based on MCTS<sup>13–15,36</sup>.

MCTS has previously been combined with a policy that is used to narrow the beam of the search tree to high-probability moves<sup>13</sup>; or to bias the bonus term towards high-probability moves<sup>48</sup>. MCTS has also been combined with a value function that is used to initialize action values in newly expanded nodes<sup>16</sup>, or to mix Monte Carlo evaluation with minimax evaluation<sup>49</sup>. By contrast, AlphaGo's use of value functions is based on truncated Monte Carlo search algorithms<sup>8,9</sup>, which terminate rollouts before the end of the game and use a value function in place of the terminal reward. AlphaGo's position evaluation mixes full rollouts with truncated rollouts, resembling in some respects the well-known temporal-difference learning algorithm TD( $\lambda$ ). AlphaGo also differs from prior work by using slower but more powerful representations of the policy and value function; evaluating deep neural networks is several orders of magnitude slower than linear representations and must therefore occur asynchronously.

The performance of MCTS is to a large degree determined by the quality of the rollout policy. Prior work has focused on handcrafted patterns<sup>50</sup> or learning rollout policies by supervised learning<sup>13</sup>, reinforcement learning<sup>16</sup>, simulation balancing<sup>51,52</sup> or online adaptation<sup>30,53</sup>; however, it is known that rollout-based position evaluation is frequently inaccurate<sup>54</sup>. AlphaGo uses relatively simple rollouts, and instead addresses the challenging problem of position evaluation more directly using value networks.

**Search algorithm.** To efficiently integrate large neural networks into AlphaGo, we implemented an asynchronous policy and value MCTS algorithm (APV-MCTS). Each node  $s$  in the search tree contains edges  $(s, a)$  for all legal actions  $a \in \mathcal{A}(s)$ . Each edge stores a set of statistics,

$$\{P(s, a), N_v(s, a), N_r(s, a), W_v(s, a), W_r(s, a), Q(s, a)\}$$

where  $P(s, a)$  is the prior probability,  $W_v(s, a)$  and  $W_r(s, a)$  are Monte Carlo estimates of total action value, accumulated over  $N_v(s, a)$  and  $N_r(s, a)$  leaf evaluations and rollout rewards, respectively, and  $Q(s, a)$  is the combined mean action value for that edge. Multiple simulations are executed in parallel on separate search threads. The APV-MCTS algorithm proceeds in the four stages outlined in Fig. 3.

**Selection** (Fig. 3a). The first in-tree phase of each simulation begins at the root of the search tree and finishes when the simulation reaches a leaf node at time step  $L$ . At each of these time steps,  $t < L$ , an action is selected according to the statistics in the search tree,  $a_t = \arg\max_a (Q(s_t, a) + u(s_t, a))$  using a variant of the PUCT

algorithm<sup>48</sup>,  $u(s, a) = c_{\text{puct}} P(s, a) \frac{\sqrt{\sum_b N_r(s, b)}}{1 + N_r(s, a)}$ , where  $c_{\text{puct}}$  is a constant determining the level of exploration; this search control strategy initially prefers actions with high prior probability and low visit count, but asymptotically prefers actions with high action value.

**Evaluation** (Fig. 3c). The leaf position  $s_L$  is added to a queue for evaluation  $v_\theta(s_L)$  by the value network, unless it has previously been evaluated. The second rollout phase of each simulation begins at leaf node  $s_L$  and continues until the end of the game. At each of these time-steps,  $t \geq L$ , actions are selected by both players according to the rollout policy,  $a_t \sim p_\pi(\cdot | s_t)$ . When the game reaches a terminal state, the outcome  $z_t = \pm r(s_T)$  is computed from the final score.

**Backup** (Fig. 3d). At each in-tree step  $t \leq L$  of the simulation, the rollout statistics are updated as if it has lost  $n_{v_l}$  games,  $N_r(s_t, a_t) \leftarrow N_r(s_t, a_t) + n_{v_l}$ ;  $W_r(s_t, a_t) \leftarrow W_r(s_t, a_t) - n_{v_l}$ ; this virtual loss<sup>55</sup> discourages other threads from simultaneously exploring the identical variation. At the end of the simulation, the rollout statistics are updated in a backward pass through each step  $t \leq L$ , replacing the virtual losses by the outcome,  $N_r(s_t, a_t) \leftarrow N_r(s_t, a_t) - n_{v_l} + 1$ ;  $W_r(s_t, a_t) \leftarrow W_r(s_t, a_t) + n_{v_l} + z_t$ . Asynchronously, a separate backward pass is initiated when the evaluation of the leaf position  $s_L$  completes. The output of the value network  $v_\theta(s_L)$  is used to update value statistics in a second backward pass through each step  $t \leq L$ ,  $N_v(s_t, a_t) \leftarrow N_v(s_t, a_t) + 1$ ,  $W_v(s_t, a_t) \leftarrow W_v(s_t, a_t) + v_\theta(s_L)$ . The overall evaluation of each state action is a weighted average of the Monte Carlo estimates,  $Q(s, a) = (1 - \lambda) \frac{W_v(s, a)}{N_v(s, a)} + \lambda \frac{W_r(s, a)}{N_r(s, a)}$ , that mixes together the value network and rollout evaluations with weighting parameter  $\lambda$ . All updates are performed lock-free<sup>56</sup>.

**Expansion** (Fig. 3b). When the visit count exceeds a threshold,  $N_r(s, a) > n_{\text{thr}}$ , the successor state  $s' = f(s, a)$  is added to the search tree. The new node is initialized to  $\{N(s', a) = N_r(s', a) = 0, W(s', a) = W_r(s', a) = 0, P(s', a) = p_\pi(a|s')\}$ , using a tree policy  $p_\pi(a|s')$  (similar to the rollout policy but with more features, see Extended Data Table 4) to provide placeholder prior probabilities for action selection. The position  $s'$  is also inserted into a queue for asynchronous GPU evaluation by the policy network. Prior probabilities are computed by the SL policy network  $p_\pi^\beta(\cdot | s')$  with a softmax temperature set to  $\beta$ ; these replace the placeholder prior probabilities,  $P(s', a) \leftarrow p_\pi^\beta(a|s')$ , using an atomic update. The threshold  $n_{\text{thr}}$  is adjusted dynamically to ensure that the rate at which positions are added to the policy queue matches the rate at which the GPUs evaluate the policy network. Positions are evaluated by both the policy network and the value network using a mini-batch size of 1 to minimize end-to-end evaluation time.

We also implemented a distributed APV-MCTS algorithm. This architecture consists of a single master machine that executes the main search, many remote worker CPUs that execute asynchronous rollouts, and many remote worker GPUs that execute asynchronous policy and value network evaluations. The entire search tree is stored on the master, which only executes the in-tree phase of each simulation. The leaf positions are communicated to the worker CPUs, which execute the rollout phase of simulation, and to the worker GPUs, which compute network features and evaluate the policy and value networks. The prior probabilities of the policy network are returned to the master, where they replace placeholder prior probabilities at the newly expanded node. The rewards from rollouts and the value network outputs are each returned to the master, and backed up the originating search path.

At the end of search AlphaGo selects the action with maximum visit count; this is less sensitive to outliers than maximizing action value<sup>15</sup>. The search tree is reused at subsequent time steps: the child node corresponding to the played action becomes the new root node; the subtree below this child is retained along with all its statistics, while the remainder of the tree is discarded. The match version of AlphaGo continues searching during the opponent's move. It extends the search

if the action maximizing visit count and the action maximizing action value disagree. Time controls were otherwise shaped to use most time in the middle-game<sup>57</sup>. AlphaGo resigns when its overall evaluation drops below an estimated 10% probability of winning the game, that is,  $\max_a Q(s, a) < -0.8$ .

AlphaGo does not employ the all-moves-as-first<sup>10</sup> or rapid action value estimation<sup>58</sup> heuristics used in the majority of Monte Carlo Go programs; when using policy networks as prior knowledge, these biased heuristics do not appear to give any additional benefit. In addition AlphaGo does not use progressive widening<sup>13</sup>, dynamic *komi*<sup>59</sup> or an opening book<sup>60</sup>. The parameters used by AlphaGo in the Fan Hui match are listed in Extended Data Table 5.

**Rollout policy.** The rollout policy  $p_\pi(a|s)$  is a linear softmax policy based on fast, incrementally computed, local pattern-based features consisting of both 'response' patterns around the previous move that led to state  $s$ , and 'non-response' patterns around the candidate move  $a$  in state  $s$ . Each non-response pattern is a binary feature matching a specific  $3 \times 3$  pattern centred on  $a$ , defined by the colour (black, white, empty) and liberty count (1, 2,  $\geq 3$ ) for each adjacent intersection. Each response pattern is a binary feature matching the colour and liberty count in a 12-point diamond-shaped pattern<sup>21</sup> centred around the previous move. Additionally, a small number of handcrafted local features encode common-sense Go rules (see Extended Data Table 4). Similar to the policy network, the weights  $\pi$  of the rollout policy are trained from 8 million positions from human games on the Tygem server to maximize log likelihood by stochastic gradient descent. Rollouts execute at approximately 1,000 simulations per second per CPU thread on an empty board.

Our rollout policy  $p_\pi(a|s)$  contains less handcrafted knowledge than state-of-the-art Go programs<sup>13</sup>. Instead, we exploit the higher-quality action selection within MCTS, which is informed both by the search tree and the policy network. We introduce a new technique that caches all moves from the search tree and then plays similar moves during rollouts; a generalization of the 'last good reply' heuristic<sup>53</sup>. At every step of the tree traversal, the most probable action is inserted into a hash table, along with the  $3 \times 3$  pattern context (colour, liberty and stone counts) around both the previous move and the current move. At each step of the rollout, the pattern context is matched against the hash table; if a match is found then the stored move is played with high probability.

**Symmetries.** In previous work, the symmetries of Go have been exploited by using rotationally and reflectionally invariant filters in the convolutional layers<sup>24,28,29</sup>. Although this may be effective in small neural networks, it actually hurts performance in larger networks, as it prevents the intermediate filters from identifying specific asymmetric patterns<sup>23</sup>. Instead, we exploit symmetries at run-time by dynamically transforming each position  $s$  using the dihedral group of eight reflections and rotations,  $d_1(s), \dots, d_8(s)$ . In an explicit symmetry ensemble, a mini-batch of all 8 positions is passed into the policy network or value network and computed in parallel. For the value network, the output values are simply averaged,  $\bar{v}_\theta(s) = \frac{1}{8} \sum_{j=1}^8 v_\theta(d_j(s))$ . For the policy network, the planes of output probabilities are rotated/reflected back into the original orientation, and averaged together to provide an ensemble prediction,  $\bar{p}_\pi(\cdot|s) = \frac{1}{8} \sum_{j=1}^8 d_j^{-1}(p_\pi(\cdot|d_j(s)))$ ; this approach was used in our raw network evaluation (see Extended Data Table 3). Instead, APV-MCTS makes use of an implicit symmetry ensemble that randomly selects a single rotation/reflection  $j \in [1, 8]$  for each evaluation. We compute exactly one evaluation for that orientation only; in each simulation we compute the value of leaf node  $s_L$  by  $v_\theta(d_j(s_L))$ , and allow the search procedure to average over these evaluations. Similarly, we compute the policy network for a single, randomly selected rotation/reflection,  $d_j^{-1}(p_\pi(\cdot|d_j(s)))$ .

**Policy network: classification.** We trained the policy network  $p_\pi$  to classify positions according to expert moves played in the KGS data set. This data set contains 29.4 million positions from 160,000 games played by KGS 6 to 9 *dan* human players; 35.4% of the games are handicap games. The data set was split into a test set (the first million positions) and a training set (the remaining 28.4 million positions). Pass moves were excluded from the data set. Each position consisted of a raw board description  $s$  and the move  $a$  selected by the human. We augmented the data set to include all eight reflections and rotations of each position. Symmetry augmentation and input features were pre-computed for each position. For each training step, we sampled a randomly selected mini-batch of  $m$  samples from the augmented KGS data set,  $\{s^k, a^k\}_{k=1}^m$  and applied an asynchronous stochastic gradient descent update to maximize the log likelihood of the action,  $\Delta\sigma = \frac{\alpha}{m} \sum_{k=1}^m \frac{\partial \log p_\pi(a^k|s^k)}{\partial \sigma}$ . The step size  $\alpha$  was initialized to 0.003 and was halved every 80 million training steps, without momentum terms, and a mini-batch size of  $m = 16$ . Updates were applied asynchronously on 50 GPUs using DistBelief<sup>61</sup>; gradients older than 100 steps were discarded. Training took around 3 weeks for 340 million training steps.

**Policy network: reinforcement learning.** We further trained the policy network by policy gradient reinforcement learning<sup>25,26</sup>. Each iteration consisted of a mini-batch of  $n$  games played in parallel, between the current policy network  $p_\rho$  that is being trained, and an opponent  $p_\rho^-$  that uses parameters  $\rho^-$  from a previous iteration, randomly sampled from a pool of opponents, so as to increase the stability of training. Weights were initialized to  $\rho = \rho^- = \sigma$ . Every 500 iterations, we added the current parameters  $\rho$  to the opponent pool. Each game  $i$  in the mini-batch was played out until termination at step  $T$ , and then scored to determine the outcome  $z_i = \pm r(s_T)$  from each player's perspective. The games were then replayed to determine the policy gradient update,  $\Delta\rho = \frac{\alpha}{n} \sum_{i=1}^n \sum_{t=1}^T \frac{\partial \log p_\rho(a_t^i|s_t^i)}{\partial \rho} (z_i - v(s_t^i))$ , using the REINFORCE algorithm<sup>25</sup> with baseline  $v(s_t^i)$  for variance reduction. On the first pass through the training pipeline, the baseline was set to zero; on the second pass we used the value network  $v_\theta(s)$  as a baseline; this provided a small performance boost. The policy network was trained in this way for 10,000 mini-batches of 128 games, using 50 GPUs, for one day. 去拟合, 去接近

**Value network: regression.** We trained a value network  $v_\theta(s) \approx v_{p_\rho}(s)$  to approximate the value function of the RL policy network  $p_\rho$ . To avoid overfitting to the strongly correlated positions within games, we constructed a new data set of uncorrelated self-play positions. This data set consisted of over 30 million positions, each drawn from a unique game of self-play. Each game was generated in three phases by randomly sampling a time step  $U \sim \text{unif}\{1, 450\}$ , and sampling the first  $t = 1, \dots, U-1$  moves from the SL policy network,  $a_t \sim p_\sigma(\cdot|s_t)$ ; then sampling one move uniformly at random from available moves,  $a_U \sim \text{unif}\{1, 361\}$  (repeatedly until  $a_U$  is legal); then sampling the remaining sequence of moves until the game terminates,  $t = U+1, \dots, T$ , from the RL policy network,  $a_t \sim p_\rho(\cdot|s_t)$ . Finally, the game is scored to determine the outcome  $z_t = \pm r(s_T)$ . Only a single training example  $(s_{U+1}, z_{U+1})$  is added to the data set from each game. This data provides unbiased samples of the value function  $v_{p_\rho}(s_{U+1}) = \mathbb{E}[z_{U+1}|s_{U+1}, a_{U+1}, \dots, T \sim p_\rho]$ . During the first two phases of generation we sample from noisier distributions so as to increase the diversity of the data set. The training method was identical to SL policy network training, except that the parameter update was based on

mean squared error between the predicted values and the observed rewards,  $\Delta\theta = \frac{\alpha}{m} \sum_{k=1}^m (z^k - v_\theta(s^k)) \frac{\partial v_\theta(s^k)}{\partial \theta}$ . The value network was trained for 50 million mini-batches of 32 positions, using 50 GPUs, for one week.

**Features for policy/value network.** Each position  $s$  was pre-processed into a set of  $19 \times 19$  feature planes. The features that we use come directly from the raw representation of the game rules, indicating the status of each intersection of the Go board: stone colour, liberties (adjacent empty points of stone's chain), captures, legality, turns since stone was played, and (for the value network only) the current colour to play. In addition, we use one simple tactical feature that computes the outcome of a ladder search<sup>7</sup>. All features were computed relative to the current colour to play; for example, the stone colour at each intersection was represented as either player or opponent rather than black or white. Each integer feature value is split into multiple  $19 \times 19$  planes of binary values (one-hot encoding). For example, separate binary feature planes are used to represent whether an intersection has 1 liberty, 2 liberties,  $\dots, \geq 8$  liberties. The full set of feature planes are listed in Extended Data Table 2.

**Neural network architecture.** The input to the policy network is a  $19 \times 19 \times 48$  image stack consisting of 48 feature planes. The first hidden layer zero pads the input into a  $23 \times 23$  image, then convolves  $k$  filters of kernel size  $5 \times 5$  with stride 1 with the input image and applies a rectifier nonlinearity. Each of the subsequent hidden layers 2 to 12 zero pads the respective previous hidden layer into a  $21 \times 21$  image, then convolves  $k$  filters of kernel size  $3 \times 3$  with stride 1, again followed by a rectifier nonlinearity. The final layer convolves 1 filter of kernel size  $1 \times 1$  with stride 1, with a different bias for each position, and applies a softmax function. The match version of AlphaGo used  $k = 192$  filters; Fig. 2b and Extended Data Table 3 additionally show the results of training with  $k = 128, 256$  and 384 filters.

The input to the value network is also a  $19 \times 19 \times 48$  image stack, with an additional binary feature plane describing the current colour to play. Hidden layers 2 to 11 are identical to the policy network, hidden layer 12 is an additional convolution layer, hidden layer 13 convolves 1 filter of kernel size  $1 \times 1$  with stride 1, and hidden layer 14 is a fully connected linear layer with 256 rectifier units. The output layer is a fully connected linear layer with a single tanh unit.

**Evaluation.** We evaluated the relative strength of computer Go programs by running an internal tournament and measuring the Elo rating of each program. We estimate the probability that program  $a$  will beat program  $b$  by a logistic function  $p(a \text{ beats } b) = \frac{1}{1 + \exp(c_{\text{elo}}(e(b) - e(a)))}$ , and estimate the ratings  $e(\cdot)$  by Bayesian logistic regression, computed by the *BayesElo* program<sup>37</sup> using the standard constant  $c_{\text{elo}} = 1/400$ . The scale was anchored to the *BayesElo* rating of professional



Go player Fan Hui (2,908 at date of submission)<sup>62</sup>. All programs received a maximum of 5 s computation time per move; games were scored using Chinese rules with a *komi* of 7.5 points (extra points to compensate white for playing second). We also played handicap games where AlphaGo played white against existing Go programs; for these games we used a non-standard handicap system in which *komi* was retained but black was given additional stones on the usual handicap points. Using these rules, a handicap of  $K$  stones is equivalent to giving  $K - 1$  free moves to black, rather than  $K - 1/2$  free moves using standard no-*komi* handicap rules. We used these handicap rules because AlphaGo's value network was trained specifically to use a *komi* of 7.5.

With the exception of distributed AlphaGo, each computer Go program was executed on its own single machine, with identical specifications, using the latest available version and the best hardware configuration supported by that program (see Extended Data Table 6). In Fig. 4, approximate ranks of computer programs are based on the highest KGS rank achieved by that program; however, the KGS version may differ from the publicly available version.

The match against Fan Hui was arbitrated by an impartial referee. Five formal games and five informal games were played with 7.5 *komi*, no handicap, and Chinese rules. AlphaGo won these games 5–0 and 3–2 respectively (Fig. 6 and Extended Data Table 1). Time controls for formal games were 1 h main time plus three periods of 30 s *byoyomi*. Time controls for informal games were three periods of 30 s *byoyomi*. Time controls and playing conditions were chosen by Fan Hui in advance of the match; it was also agreed that the overall match outcome would be determined solely by the formal games. To approximately assess the relative rating of Fan Hui to computer Go programs, we appended the results of all ten games to our internal tournament results, ignoring differences in time controls.

39. Littman, M. L. Markov games as a framework for multi-agent reinforcement learning. In *11th International Conference on Machine Learning*, 157–163 (1994).
40. Knuth, D. E. & Moore, R. W. An analysis of alpha-beta pruning. *Artif. Intell.* **6**, 293–326 (1975).
41. Sutton, R. Learning to predict by the method of temporal differences. *Mach. Learn.* **3**, 9–44 (1988).
42. Baxter, J., Tridgell, A. & Weaver, L. Learning to play chess using temporal differences. *Mach. Learn.* **40**, 243–263 (2000).
43. Veness, J., Silver, D., Blair, A. & Uther, W. Bootstrapping from game tree search. In *Advances in Neural Information Processing Systems* (2009).
44. Samuel, A. L. Some studies in machine learning using the game of checkers II - recent progress. *IBM J. Res. Develop.* **11**, 601–617 (1967).
45. Schaeffer, J., Hlynka, M. & Jussila, V. Temporal difference learning applied to a high-performance game-playing program. In *17th International Joint Conference on Artificial Intelligence*, 529–534 (2001).
46. Tesauro, G. TD-gammon, a self-teaching backgammon program, achieves master-level play. *Neural Comput.* **6**, 215–219 (1994).
47. Dahl, F. Honte, a Go-playing program using neural nets. In *Machines that learn to play games*, 205–223 (Nova Science, 1999).
48. Rosin, C. D. Multi-armed bandits with episode context. *Ann. Math. Artif. Intell.* **61**, 203–230 (2011).
49. Lanctot, M., Winands, M. H. M., Pepels, T. & Sturtevant, N. R. Monte Carlo tree search with heuristic evaluations using implicit minimax backups. In *IEEE Conference on Computational Intelligence and Games*, 1–8 (2014).
50. Gelly, S., Wang, Y., Munos, R. & Teytaud, O. Modification of UCT with patterns in Monte-Carlo Go. Tech. Rep. 6062, INRIA (2006).
51. Silver, D. & Tesauro, G. Monte-Carlo simulation balancing. In *26th International Conference on Machine Learning*, 119 (2009).
52. Huang, S.-C., Coulom, R. & Lin, S.-S. Monte-Carlo simulation balancing in practice. In *7th International Conference on Computers and Games*, 81–92 (Springer-Verlag, 2011).
53. Baier, H. & Drake, P. D. The power of forgetting: improving the last-good-reply policy in Monte Carlo Go. *IEEE Trans. Comput. Intell. AI in Games* **2**, 303–309 (2010).
54. Huang, S. & Müller, M. Investigating the limits of Monte-Carlo tree search methods in computer Go. In *8th International Conference on Computers and Games*, 39–48 (2013).
55. Segal, R. B. On the scalability of parallel UCT. *Computers and Games* **6515**, 36–47 (2011).
56. Enzenberger, M. & Müller, M. A lock-free multithreaded Monte-Carlo tree search algorithm. In *12th Advances in Computer Games Conference*, 14–20 (2009).
57. Huang, S.-C., Coulom, R. & Lin, S.-S. Time management for Monte-Carlo tree search applied to the game of Go. In *International Conference on Technologies and Applications of Artificial Intelligence*, 462–466 (2010).
58. Gelly, S. & Silver, D. Monte-Carlo tree search and rapid action value estimation in computer Go. *Artif. Intell.* **175**, 1856–1875 (2011).
59. Baudiš, P. Balancing MCTS by dynamically adjusting the *komi* value. *ICGA J.* **34**, 131 (2011).
60. Baier, H. & Winands, M. H. Active opening book application for Monte-Carlo tree search in 19×19 Go. In *Benelux Conference on Artificial Intelligence*, 3–10 (2011).
61. Dean, J. et al. Large scale distributed deep networks. In *Advances in Neural Information Processing Systems*, 1223–1231 (2012).
62. Go ratings. <http://www.goratings.org>.

**Extended Data Table 1 | Details of match between AlphaGo and Fan Hui**

Date	Black	White	Category	Result
5/10/15	Fan Hui	<i>AlphaGo</i>	Formal	<i>AlphaGo</i> wins by 2.5 points
5/10/15	Fan Hui	<i>AlphaGo</i>	Informal	Fan Hui wins by resignation
6/10/15	<i>AlphaGo</i>	Fan Hui	Formal	<i>AlphaGo</i> wins by resignation
6/10/15	<i>AlphaGo</i>	Fan Hui	Informal	<i>AlphaGo</i> wins by resignation
7/10/15	Fan Hui	<i>AlphaGo</i>	Formal	<i>AlphaGo</i> wins by resignation
7/10/15	Fan Hui	<i>AlphaGo</i>	Informal	<i>AlphaGo</i> wins by resignation
8/10/15	<i>AlphaGo</i>	Fan Hui	Formal	<i>AlphaGo</i> wins by resignation
8/10/15	<i>AlphaGo</i>	Fan Hui	Informal	<i>AlphaGo</i> wins by resignation
9/10/15	Fan Hui	<i>AlphaGo</i>	Formal	<i>AlphaGo</i> wins by resignation
9/10/15	<i>AlphaGo</i>	Fan Hui	Informal	Fan Hui wins by resignation

The match consisted of five formal games with longer time controls, and five informal games with shorter time controls. Time controls and playing conditions were chosen by Fan Hui in advance of the match.

Extended Data Table 2 | Input features for neural networks

Feature	# of planes	Description
Stone colour	3	Player stone / opponent stone / empty
Ones	1	A constant plane filled with 1
Turns since	8	How many turns since a move was played
Liberties	8	Number of liberties (empty adjacent points)
Capture size	8	How many opponent stones would be captured
Self-atari size	8	How many of own stones would be captured
Liberties after move	8	Number of liberties after this move is played
Ladder capture	1	Whether a move at this point is a successful ladder capture
Ladder escape	1	Whether a move at this point is a successful ladder escape
Sensibleness	1	Whether a move is legal and does not fill its own eyes
Zeros	1	A constant plane filled with 0
Player color	1	Whether current player is black

Feature planes used by the policy network (all but last feature) and value network (all features).



Extended Data Table 3 | Supervised learning results for the policy network

Architecture			Evaluation					
Filters	Symmetries	Features	Test accu- racy %	Train accu- racy %	Raw net wins %	<i>AlphaGo</i> wins %	Forward time (ms)	
128	1	48	54.6	57.0	36	53	2.8	
192	1	48	55.4	58.0	50	50	4.8	
256	1	48	55.9	59.1	67	55	7.1	
256	2	48	56.5	59.8	67	38	13.9	
256	4	48	56.9	60.2	69	14	27.6	
256	8	48	57.0	60.4	69	5	55.3	
192	1	4	47.6	51.4	25	15	4.8	
192	1	12	54.7	57.1	30	34	4.8	
192	1	20	54.7	57.2	38	40	4.8	
192	8	4	49.2	53.2	24	2	36.8	
192	8	12	55.7	58.3	32	3	36.8	
192	8	20	55.8	58.4	42	3	36.8	

The policy network architecture consists of 128, 192 or 256 filters in convolutional layers; an explicit symmetry ensemble over 2, 4 or 8 symmetries; using only the first 4, 12 or 20 input feature planes listed in Extended Data Table 1. The results consist of the test and train accuracy on the KGS data set; and the percentage of games won by given policy network against AlphaGo's policy network (highlighted row 2): using the policy networks to select moves directly (raw wins); or using AlphaGo's search to select moves (AlphaGo wins); and finally the computation time for a single evaluation of the policy network.

**Extended Data Table 4 | Input features for rollout and tree policy**

Feature	# of patterns	Description
Response	1	Whether move matches one or more response pattern features
Save atari	1	Move saves stone(s) from capture
Neighbour	8	Move is 8-connected to previous move
Nakade	8192	Move matches a <i>nakade</i> pattern at captured stone
Response pattern	32207	Move matches 12-point diamond pattern near previous move
Non-response pattern	69338	Move matches $3 \times 3$ pattern around move
Self-atari	1	Move allows stones to be captured
Last move distance	34	Manhattan distance to previous two moves
Non-response pattern	32207	Move matches 12-point diamond pattern centred around move

Features used by the rollout policy (first set) and tree policy (first and second set). Patterns are based on stone colour (black/white/empty) and liberties (1, 2,  $\geq 3$ ) at each intersection of the pattern.

**Extended Data Table 5 | Parameters used by AlphaGo**

Symbol	Parameter	Value
$\beta$	Softmax temperature	0.67
$\lambda$	Mixing parameter	0.5
$n_{\text{vl}}$	Virtual loss	3
$n_{\text{thr}}$	Expansion threshold	40
$c_{\text{puct}}$	Exploration constant	5



Extended Data Table 6 | Results of a tournament between different Go programs

Short name	Computer Player	Version	Time settings	CPUs	GPUs	KGS Rank	Elo
$\alpha_{rvp}^d$	Distributed AlphaGo	See Methods	5 seconds	1202	176	–	3140
$\alpha_{rvp}$	AlphaGo	See Methods	5 seconds	48	8	–	2890
<i>CS</i>	CrazyStone	2015	5 seconds	32	–	6d	1929
<i>ZN</i>	Zen	5	5 seconds	8	–	6d	1888
<i>PC</i>	Pachi	10.99	400,000 sims	16	–	2d	1298
<i>FG</i>	Fuego	svn1989	100,000 sims	16	–	–	1148
<i>GG</i>	GnuGo	3.8	level 10	1	–	5k	431
$CS_4$	CrazyStone	4 handicap stones	5 seconds	32	–	–	2526
$ZN_4$	Zen	4 handicap stones	5 seconds	8	–	–	2413
$PC_4$	Pachi	4 handicap stones	400,000 sims	16	–	–	1756

Each program played with a maximum of 5 s thinking time per move; the games against Fan Hui were conducted using longer time controls, as described in Methods.  $CN_4$ ,  $ZN_4$  and  $PC_4$  were given 4 handicap stones; *komi* was 7.5 in all games. Elo ratings were computed by BayesElo.

Extended Data Table 7 | Results of a tournament between different variants of AlphaGo

Short name	Policy network	Value network	Rollouts	Mixing constant	Policy GPUs	Value GPUs	Elo rating
$\alpha_{rvp}$	$p_\sigma$	$v_\theta$	$p_\pi$	$\lambda = 0.5$	2	6	2890
$\alpha_{vp}$	$p_\sigma$	$v_\theta$	—	$\lambda = 0$	2	6	2177
$\alpha_{rp}$	$p_\sigma$	—	$p_\pi$	$\lambda = 1$	8	0	2416
$\alpha_{rv}$	$[p_\tau]$	$v_\theta$	$p_\pi$	$\lambda = 0.5$	0	8	2077
$\alpha_v$	$[p_\tau]$	$v_\theta$	—	$\lambda = 0$	0	8	1655
$\alpha_r$	$[p_\tau]$	—	$p_\pi$	$\lambda = 1$	0	0	1457
$\alpha_p$	$p_\sigma$	—	—	—	0	0	1517

Evaluating positions using rollouts only ( $\alpha_{rp}$ ,  $\alpha_r$ ), value nets only ( $\alpha_{vp}$ ,  $\alpha_v$ ), or mixing both ( $\alpha_{rvp}$ ,  $\alpha_{rv}$ ); either using the policy network  $p_\sigma(\alpha_{rvp}$ ,  $\alpha_{vp}$ ,  $\alpha_{rp}$ ), or no policy network ( $\alpha_{rvp}$ ,  $\alpha_{vp}$ ,  $\alpha_{rp}$ ), that is, instead using the placeholder probabilities from the tree policy  $p_\tau$  throughout. Each program used 5 s per move on a single machine with 48 CPUs and 8 GPUs. Elo ratings were computed by BayesElo.

**Extended Data Table 8 | Results of a tournament between AlphaGo and distributed AlphaGo, testing scalability with hardware**

<i>AlphaGo</i>	Search threads	CPUs	GPUs	Elo
Asynchronous	1	48	8	2203
Asynchronous	2	48	8	2393
Asynchronous	4	48	8	2564
Asynchronous	8	48	8	2665
Asynchronous	16	48	8	2778
Asynchronous	32	48	8	2867
Asynchronous	40	48	8	2890
Asynchronous	40	48	1	2181
Asynchronous	40	48	2	2738
Asynchronous	40	48	4	2850
Distributed	12	428	64	2937
Distributed	24	764	112	3079
Distributed	40	1202	176	3140
Distributed	64	1920	280	3168

Each program played with a maximum of 2 s thinking time per move. Elo ratings were computed by BayesElo.



Extended Data Table 9 | Cross-table of win rates in per cent between programs

	$\alpha_{rvp}$	$\alpha_{vp}$	$\alpha_{rp}$	$\alpha_{rv}$	$\alpha_r$	$\alpha_v$	$\alpha_p$
$\alpha_{rvp}$	-	1 [0; 5]	5 [4; 7]	0 [0; 4]	0 [0; 8]	0 [0; 19]	0 [0; 19]
$\alpha_{vp}$	99 [95; 100]	-	61 [52; 69]	35 [25; 48]	6 [1; 27]	0 [0; 22]	1 [0; 6]
$\alpha_{rp}$	95 [93; 96]	39 [31; 48]	-	13 [7; 23]	0 [0; 9]	0 [0; 22]	4 [1; 21]
$\alpha_{rv}$	100 [96; 100]	65 [52; 75]	87 [77; 93]	-	0 [0; 18]	29 [8; 64]	48 [33; 65]
$\alpha_r$	100 [92; 100]	94 [73; 99]	100 [91; 100]	100 [82; 100]	-	78 [45; 94]	78 [71; 84]
$\alpha_v$	100 [81; 100]	100 [78; 100]	100 [78; 100]	71 [36; 92]	22 [6; 55]	-	30 [16; 48]
$\alpha_p$	100 [81; 100]	99 [94; 100]	96 [79; 99]	52 [35; 67]	22 [16; 29]	70 [52; 84]	-
$CS$	100 [97; 100]	74 [66; 81]	98 [94; 99]	80 [70; 87]	5 [3; 7]	36 [16; 61]	8 [5; 14]
$ZN$	99 [93; 100]	84 [67; 93]	98 [93; 99]	92 [67; 99]	6 [2; 19]	40 [12; 77]	100 [65; 100]
$PC$	100 [98; 100]	99 [95; 100]	100 [98; 100]	98 [89; 100]	78 [73; 81]	87 [68; 95]	55 [47; 62]
$FG$	100 [97; 100]	99 [93; 100]	100 [96; 100]	100 [91; 100]	78 [73; 83]	100 [65; 100]	65 [55; 73]
$GG$	100 [44; 100]	100 [34; 100]	100 [68; 100]	100 [57; 100]	99 [97; 100]	67 [21; 94]	99 [95; 100]
$CS_4$	77 [69; 84]	12 [8; 18]	53 [44; 61]	15 [8; 24]	0 [0; 3]	0 [0; 30]	0 [0; 8]
$ZN_4$	86 [77; 92]	25 [16; 38]	67 [56; 76]	14 [7; 27]	0 [0; 12]	0 [0; 43]	-
$PC_4$	99 [97; 100]	82 [75; 88]	98 [95; 99]	89 [79; 95]	32 [26; 39]	13 [3; 36]	35 [25; 46]

95% Agresti–Coull confidence intervals in grey. Each program played with a maximum of 5 s thinking time per move.  $CN_4$ ,  $ZN_4$  and  $PC_4$  were given 4 handicap stones; *komi* was 7.5 in all games. Distributed AlphaGo scored 77% [70; 82] against  $\alpha_{vp}$  and 100% against all other programs (no handicap games were played).

Extended Data Table 10 | Cross-table of win rates in per cent between programs in the single-machine scalability study

Threads	GPU	1	2	4	8	16	32	40	40	40	40
		8	8	8	8	8	8	8	4	2	1
1	8	-	70 [61;78]	90 [84;94]	94 [83;98]	86 [72;94]	98 [91;100]	98 [92;99]	100 [76;100]	96 [91;98]	38 [25;52]
2	8	30 [22;39]	-	72 [61;81]	81 [71;88]	86 [76;93]	92 [83;97]	93 [86;96]	83 [69;91]	84 [75;90]	26 [17;38]
4	8	10 [6;16]	28 [19;39]	-	62 [53;70]	71 [61;80]	82 [71;89]	84 [74;90]	81 [69;89]	78 [63;88]	18 [10;28]
8	8	6 [2;17]	19 [12;29]	38 [30;47]	-	61 [51;71]	65 [51;76]	73 [62;82]	74 [59;85]	64 [55;73]	12 [3;34]
16	8	14 [6;28]	14 [7;24]	29 [20;39]	39 [29;49]	-	52 [41;63]	61 [50;71]	52 [41;64]	41 [32;51]	5 [1;25]
32	8	2 [0;9]	8 [3;17]	18 [11;29]	35 [24;49]	48 [37;59]	-	52 [42;63]	44 [32;57]	26 [17;36]	0 [0;30]
40	8	2 [1;8]	8 [4;14]	16 [10;26]	27 [18;38]	39 [29;50]	48 [37;58]	-	43 [30;56]	41 [26;58]	4 [1;18]
40	4	0 [0;24]	17 [9;31]	19 [11;31]	26 [15;41]	48 [36;59]	56 [43;68]	57 [44;70]	-	29 [18;41]	2 [0;11]
40	2	4 [2;9]	16 [10;25]	22 [12;37]	36 [27;45]	59 [49;68]	74 [64;83]	59 [42;74]	71 [59;82]	-	5 [1;17]
40	1	62 [48;75]	74 [62;83]	82 [72;90]	88 [66;97]	95 [75;99]	100 [70;100]	96 [82;99]	98 [89;100]	95 [83;99]	-

95% Agresti–Coull confidence intervals in grey. Each program played with 2 s per move; *komi* was 7.5 in all games.

**Extended Data Table 11 | Cross-table of win rates in per cent between programs in the distributed scalability study**

Threads			40	12	24	40	64
GPU			8	64	112	176	280
CPU			48	428	764	1202	1920
40	8	48	-	52 [43; 61]	68 [59; 76]	77 [70; 82]	81 [65; 91]
12	64	428	48 [39; 57]	-	64 [54; 73]	62 [41; 79]	83 [55; 95]
24	112	764	32 [24; 41]	36 [27; 46]	-	36 [20; 57]	60 [51; 69]
40	176	1202	23 [18; 30]	38 [21; 59]	64 [43; 80]	-	53 [39; 67]
64	280	1920	19 [9; 35]	17 [5; 45]	40 [31; 49]	47 [33; 61]	-

95% Agresti–Coull confidence intervals in grey. Each program played with 2 s per move; *komi* was 7.5 in all games.