

Self-Play Learning Without a Reward Metric

Dan Schmidt,¹ Nick Moran,¹ Jonathan S. Rosenfeld,²
Jonathan Rosenthal,³ Jonathan Yedidia¹

¹Analog Devices, ²Massachusetts Institute of Technology, ³Purdue University
{daniel.schmidt, nicholas.moran, jonathan.yedidia}@analog.com, jonsr@csail.mit.edu, rosensth0@purdue.edu

Abstract

The AlphaZero algorithm for the learning of strategy games via self-play, which has produced superhuman ability in the games of Go, chess, and shogi, uses a quantitative reward function for game outcomes, requiring the users of the algorithm to explicitly balance different components of the reward against each other, such as the game winner and margin of victory. We present a modification to the AlphaZero algorithm that requires only a total ordering over game outcomes, obviating the need to perform any quantitative balancing of reward components. We demonstrate that this system learns optimal play in a comparable amount of time to AlphaZero on a sample game.

Introduction

The AlphaZero (Silver et al. 2017) algorithm learns to master a two-player competitive game starting with no knowledge except for the rules of the game. As with any sort of reinforcement learning system, it requires a reward function so that good outcomes can be distinguished from bad ones. In the case of AlphaZero, this is by default a binary-valued function, simply distinguishing wins from losses; for chess a third intermediate value is added to represent draws. Since no differentiation is made between different sorts of wins, the learner has no explicit incentive to win more convincingly by scoring more points (in the case of games such as Go) or by concluding the game in fewer moves (in the case of games such as chess and shogi). As a result, the trained agent, although superhuman in most aspects of the game, often makes clearly suboptimal moves from a human point of view when victory is assured, since this does not affect the reward received.

In fact, in practice these slack moves can occasionally cause its advantage to slowly slip away, and the result of the game can change for the worse, indicating that focusing only on the category of result can impede generalization. For these reasons, it is desirable if possible to introduce a secondary objective (score differential or number of moves) so that the learner can play optimally from that perspective throughout the entire game.

The simplest way to add a secondary objective is to modify the reward function so that more extreme game outcomes

are associated with rewards and penalties that are greater in magnitude. However, this requires the experimenter to set hyperparameters defining an explicit quantitative trade-off between the primary and secondary objectives.

Our variant of the AlphaZero algorithm requires only a total ordering between game outcomes and produces an agent that attempts to maximize the rank of the outcomes of its games in that ordering, without the need for any explicit quantitative balancing factor. We introduce a method of learning optimal play given this ordering, as well as a network architecture for predicting outcomes from game states that is independent of reward function. Because we use only the ordering of outcomes and not any associated value, we can learn in the absence of a metric that defines an explicit real-valued distance between outcomes.

Related Work

AlphaZero (Silver et al. 2017) is the ancestor of all the learning algorithms considered here. It explicitly ignores any secondary objectives such as score differential. Open-source software designed to reproduce the AlphaZero algorithm, such as Leela Zero¹ and Leela Chess Zero², share its limitations in this respect.

KataGo (Wu 2019) attempts to improve upon AlphaZero in various respects in the specific domain of Go, including dispensing a greater reward for larger wins. The additional reward as a function of score difference is an ad hoc curve, tuned by hand, that works well in practice.

The Go engine SAI (Morandini et al. 2019) attempts to maximize the margin of victory by awarding artificial bonus points to the losing player, inducing the winning player to increase the actual score difference to overcome that handicap.

Reward shaping, the practice of engineering a reward function to improve learning performance, has a long history in the field of reinforcement learning (Ng, Harada, and Russell 1999). Through the lens of this work, our method may be viewed as a type of automatic and adaptive reward shaping.

The use of rank-based reward functions is common in the evolution strategy algorithm and its many variants (Rechen-

¹<https://zero.sjeng.org>

²<https://lczero.org>

berg 1973) (Wierstra et al. 2014). Such rank-based functions are typically used as a way to smooth the reward landscape, detach learning from hand-engineered payoffs, and provide scale-invariance of the algorithm with respect to raw reward values. Rank-based rewards have been used to adapt self-play reinforcement learning algorithms to single-player tasks, in particular combinatorial optimization problems (Laterre et al. 2018). Deep Ordinal Reinforcement Learning (Zap, Joppen, and Fürnkranz 2019) adapts Q-learning to use an ordinal reward scale (although without the use of a population ranking) to induce scale-invariance and reduce the need for manual reward-shaping.

One can avoid the need to hand-tune the magnitude of the reward function in reinforcement learning by adaptively normalizing rewards as learning progresses (van Hasselt et al. 2016).

Our method may be viewed as a form of preference-based reinforcement learning (Wirth et al. 2017), in which an agent receives feedback that indicates the relative utility of two states or actions, rather than an absolute numerical reward, removing the need for hand-tuning of reward functions.

Method

We distinguish between the raw **result** of a game (win/loss/draw) and the game’s **outcome**, which contains more detailed information, such as length of game or margin of victory, allowing us to perform a finer comparison of two games with the same result.

CDF rewards

We wish to motivate our learner to win more convincingly without having to introduce an explicit quantitative tradeoff, such as giving the winning player a reward of $1 + \gamma m$, where m is the margin of victory. Removing this explicit secondary reward function reduces the number of hyperparameters and means that we do not have to tune γ to ensure that the risk-reward ratio involved in pursuit of larger wins is optimal by some criterion.

Our approach requires only a total ordering over game outcomes. Given that ordering, we can base the reward solely upon it; at the conclusion of a game, if a player’s outcome is superior to some fraction f of a corpus of games played by similar (possibly identical) players, we give that player a reward of f .

In effect, we are calculating the cumulative distribution function of observed outcomes and using it as our reward function. Naturally, this CDF of representative outcomes changes over the course of training, so the reward corresponding to a given outcome changes as well, but the ordering of rewards continues to match the ordering of outcomes. (For clarity, we work here with a reward scale from 0 to 1, as that matches the semantics of the CDF, but for convenience our code uses a scale from -1 to $+1$ in order to turn it into a zero-sum game.)

In practice, a sliding window of some number of generations of the most recent self-play training game outcomes is maintained, from which the current CDF of outcomes is computed. Before any games are played at all, the reward function is a constant, and play is completely random.




Figure 1: CDF-based rewards at three points during training of a 3×9 opposition game agent. The first three plots indicate the distribution of recent outcomes after 5, 20, and 50 generations of self-play; the last plot indicates the resulting CDF reward functions.

The CDFs of game outcomes in a two-player zero-sum game are specific to each side and are complements of each other. In a game with a first-player advantage, a narrow win on the board may well be below the 50th percentile of outcomes for the first player, and correspondingly above the 50th percentile of outcomes for the second player. In practice we maintain a CDF from the first player’s point of view and calculate the second player’s reward accordingly.

Figure 1 illustrates the evolution of a CDF reward function over the course of training; this is the reward function induced by the “CDF/outcome” agent from Figure 3. The distribution of observed outcomes becomes more concentrated as the system learns to play more effectively. In this game, two-thirds of the starting positions result in a win for the first player with optimal play, and it can be seen that this first-player advantage is recognized by the reward function.

Outcome prediction

In the AlphaZero algorithm, one of the network outputs given a position is the expected reward of the game that is continued from that position. In our case, however, the reward corresponding to a game outcome is not fixed, so we cannot directly predict the expected value of a position, as the reward function could shift in the future, causing the value prediction to no longer match the current CDF. Instead, the network needs to predict a game outcome rather than a value.

This outcome prediction can take many different forms in practice. For example, if the number of distinct outcomes is limited, the network can simply produce a categorical output. Our experiments focus on a game whose secondary objective is to win quickly, and use a game outcome output head consisting of three categorical outputs (win, draw, or

loss) used as inputs into a softmax function, as well as two additional predictions for the number of remaining moves in the game in the separate cases of a win and a loss. To produce a value, the rewards corresponding to these three outcomes are computed with the current CDF, and a final value is computed by weighting these rewards according to the softmax output. This system captures a reasonable amount of the uncertainty in the prediction and seems to be a happy middle ground between a categorical output over a very large number of outcomes on the one hand and a single point estimate of the outcome on the other.

As an example, if the post-softmax outputs for win, loss, and draw are 0.60, 0.35, and 0.05 respectively, and the plies-remaining outputs for wins and losses are 11 and 14 after rounding, the resulting value will be $0.60 \cdot F(\text{win in } 11) + 0.35 \cdot F(\text{loss in } 14) + 0.05 \cdot F(\text{draw})$, where F is the current CDF reward function.

Note that in games where the secondary objective involves the length of the game, it is most natural for the network to predict the number of moves remaining (since this is invariant with respect to how many moves have already been played), rather than to take as input the number of moves played so far and output the total predicted game length. Once the number of moves remaining has been predicted, it can be added to the number of moves already played in order to generate a final game outcome that is comparable to other ones.

In the case of a game where we supply a larger reward for winning quickly, there is an additional benefit of outcome prediction over value prediction. A network that directly predicts the value of a position would need to take the number of moves already played as an input and take it into account when calculating a reward, whereas an outcome-based network can just predict the number of moves remaining, which can then be combined with the current game state to produce a value. This can provide an advantage over a value prediction, since the architecture automatically supplies some generalization of the reward structure.

Once an outcome, or weighted collection of likely outcomes, has been generated, it can be dynamically converted to a value using the current CDF. These values are then used in the standard manner in the remainder of the MCTS algorithm.

Training of the outcome head is performed by storing the outcome (or remaining outcome, as described above) in each training example, rather than reward, and using a suitable loss function, such as cross-entropy on the win/loss/draw softmax outputs and squared loss on the relevant remaining-move outputs. For instance, if a training example is a loss, the output corresponding to the number of remaining moves in a win has no cost associated with it.

Virtual matches and outcome bonuses

Another way of interpreting these rewards is to imagine that every game is one half of a two-game match where each player takes each side once, with a full reward given to the player with the better average game outcome over the course of the match. To increase our number of data samples, we can create a large corpus of virtual matches incorporating

every game rather than a single match per game, by pairing every individual game outcome with every other game outcome. Given a single game outcome, the expected score of the virtual matches that result from pairing it with every other game is exactly the CDF value of that game outcome.

Given this interpretation, we can explore the effects of scoring the virtual matches differently. For example, we can give the match winner 1 point if they win both individual games of the match, but only $\alpha < 1$ points if they win the match due to having a better tiebreaker. This effectively creates a discontinuity of size $1 - \alpha$ in the CDF at the boundary between losses and wins, and recognizes the fact that there is a qualitative perceived difference between the two categories of outcomes. Of course, this is an additional hyperparameter.

It is straightforward to create a new reward function that simulates these virtual matches with bonuses. Letting W and L be the number of wins and losses in the corpus, and i being the index in $[0, L + W - 1]$ indicating where the outcome in question lies in the sorted list of reference outcomes, the associated reward for a loss, when rescaled to $[-1, +1]$, is

$$-1 + \frac{L - \alpha L + 2\alpha i}{L + W - 1},$$

while the associated reward for a win is

$$1 - \frac{W - \alpha W + 2\alpha(L + W - 1 - i)}{L + W - 1},$$

and the reward for any draw is

$$\frac{(1 + \alpha)(L - W)}{2(L + W)},$$

which all reduce to the expected formulas when $\alpha = 1$.

This bonus system can be used if some important part of the reward structure is not captured by the ordering over outcomes. A feature of the CDF reward system is that it is invariant to reward scaling, but if the shape of the underlying reward function is important, we can recognize it using this sort of bonus.

Experiments

The opposition game

We use the ‘‘opposition game’’, a simple game designed to teach fundamental concepts of chess endgame strategy, as a testbed for our method because its difficulty is easily parameterizable by changing the board size, and perfect play is easily achievable, allowing us to measure the amount of training needed to reach perfection and quantitatively measure performance during training by comparison to a perfect player. In addition, from many positions there are multiple winning moves but some win faster than others, so it is a convenient testbed for training learners to win quickly. Despite its simplicity, it will be seen that a straightforward implementation of the AlphaZero algorithm has difficulty learning to play optimally from a length-of-game standpoint.

The opposition game is played on a $w \times h$ chessboard with two chess kings, each starting on its own back rank. The players alternate turns; on one’s turn one moves one’s




Figure 2: Two sample positions from the 3×9 opposition game with Black to play; White is attempting to reach the top rank while Black is attempting to reach the bottom. In the left position, the dark green move wins in 11 ply, while the light blue move wins in 15 ply; all other moves lose. In the right position, the dark green move wins, while all other moves lose.

own king by one square in any direction horizontally, vertically, or diagonally. The game is won by capturing the other player’s king (there is no concept of check) or by reaching the other player’s back rank. Two sample positions are illustrated in Figure 2. It is possible to reach a position where with perfect play neither player can make progress, but for any starting position perfect play always ends with a decisive result.

Because it is possible to reach a state where neither side can win with best play, and because a suboptimal player may not be able to find a win from a theoretically winnable position, we declare a game to be drawn if a certain number of moves have been played without a winner. In our experiments this timeout is imposed at a ply value of $20h$.

We consider all possible initial placements of the kings on their back ranks, resulting in w^2 legal starting positions. Each training game uses one of these starting positions selected uniformly at random. It may appear unprincipled to use a single CDF for multiple starting positions, since the outcome achieved with optimal play may vary widely between different starting positions; but the true goal is to work with a representative distribution of outcomes and so induce good play, not to compute a mathematically exact result. We find that using the same CDF for a combination of initial positions does not impede learning.

Results

The experimental training setup is detailed in Appendix A.

We consider four reward functions for use during training. They are all zero-sum and range from -1 to $+1$, although the extreme values of this range are not always achievable.

- The **primitive** reward function is just -1 , 0 , or $+1$, depending on the result of the game.
- The **hand-tuned** reward function is a linear interpolation

from $+1$ for a win (-1 for a loss) in zero moves to 0 for a game that has timed out.

- The **CDF** reward function is $2f - 1$, where f is the fraction of recent training games that have a worse outcome for this player’s side.
- The **CDF-bonus** reward function treats the CDF function as an expected result of virtual matches in which a match win due to a tiebreaker (such as game length) is awarded only α points rather than 1 .

The CDF reward functions require us to use a network with an outcome head rather than a value head, as the mapping from outcome to reward is dynamic. The other reward functions can be used with either a value head or an outcome head.

We evaluate the agents by pitting them against a perfect player. After every generation of training, an 18-game match is played against the perfect player, with two games played from each starting position with kings on their back ranks. A total score for the match is awarded based on cumulative game score computed with the hand-tuned reward function. This score is always nonpositive for our agents, since they are playing against the optimal strategy, so we flip the sign of the score and refer to it as being in units of demerits. An agent that lost each of the match games instantaneously (impossible in practice) would receive 18 demerits. Figure 3 illustrates the result of this evaluation.

The agents trained with the primitive reward have trouble achieving a perfect score by this measure; it takes them time to unlearn suboptimal moves that cause them to time out against the perfect player, and even at the end of training do not put up optimal resistance in losing positions, as can be seen by the fact that they do not converge to 0 demerits.

The agents trained with the hand-tuned reward reach optimal play quickly, but require exact knowledge of the specific reward structure upon which they will be eventually evaluated.

The agent trained with the CDF reward begins with completely random play, since it cannot distinguish yet between different outcomes. For this reason, it requires some additional startup time to generate a corpus of relevant outcomes, but then quickly learns to maximize its reward based purely on trying to outperform its peers without the need for a quantitative metric. The CDF-based agent that is given a winning bonus learns faster at first, although in this case it took some time to progress past a suboptimal plateau.

Discussion and Future Work

We have presented a variant of the AlphaZero algorithm that does not require specific numeric rewards to be associated with outcomes but requires only a total ordering over outcomes. This system can be used to learn games with a more interesting set of outcomes than just wins, losses and draws, without having to adjust any hyperparameters defining some specific relative importance of these additional factors.

The use of a CDF of game outcomes to define a reward function may appear problematic at first; after all, with perfect play the CDF of a nondeterministic game should have




Figure 3: Demerits over time for differently-trained agents. The legend indicates the reward function and type of network head for each agent. For the CDF-bonus agent, $\alpha = 0.5$. The plateau near 5 demerits is a state where the agent repeatedly makes moves that are theoretically winning but delay the end of the game, eventually resulting in a timeout against a perfect opponent.

all of its weight on a single outcome. But our CDF is “softened” by being assembled from training games, which are played with Dirichlet noise and a nonzero temperature (following AlphaZero), so in practice non-optimal outcomes do occur. In any case, the method requires only an approximate histogram of typical outcomes so that rough comparisons of the relative degree of their unusualness can be made. In fact in theory any monotonically increasing reward function could be used to learn optimal play, but the CDF reward function has the advantage of having a natural definition that automatically takes the rarity of different outcomes into account.

Because the CDF reward function is based on prior self-play games, play is completely random at the beginning of training and it may take some time for the games to achieve a terminal state. To jumpstart the training process, it is possible to begin training with a standard reward function and then switch to a CDF-based reward function during training. Since the network predicts outcomes, not values, it is not made obsolete when the reward function is switched out.

Our experiments have been limited to games where we would like to induce players to win faster (and lose more slowly), but it is also a natural fit for score-based games such as Go where the winner would like to maximize the margin of victory. We would like to explore games of this nature as well. Now that the ability of this method to learn has been verified in the more tractable context of the opposition game, the natural next step is to test it on more complex games that currently suffer from being trained with a binary or ternary reward function.

Appendix A: Experimental setup

We use the same basic learning method as AlphaZero (Silver et al. 2017). Each generation of training consists of 25 self-play games. After each generation of self-play games, we train our network on the most recent 5 generations of games, also using the most recent 5 generations of games to compute the CDF used for the reward function.

The network consists of three 3×3 convolutional layers with 16 channels and ReLU activations, each followed by a batch norm layer. The policy head has one more such convolutional layer with 32 channels followed by a single fully-connected layer with a softmax output. The second head consists of a fully connected layer of size 64 with ReLU activation followed by a final fully connected layer to the output(s).

When this second head is a value head, its output is sent through the tanh function to produce a value in the range $[-1, +1]$. When the second head is an outcome head, it consists of five nodes. Three of them indicate the relative probability of a win, loss, and draw result, and are fed into a softmax. Two further nodes predict the number of ply left in the game in the case of a win and a loss respectively. Therefore three distinct outcomes (a win of some length, a loss of some other length, and a draw) with varying probabilities are proposed, and their corresponding values can be weighted accordingly.

The CDF reward is determined by the last 5 generations of self-play game outcomes. The reward for outcomes that are not present in that set is calculated by interpolating between the rewards for the two recorded outcomes that bound it,

Plane	Values
0	1 at position of Player One, $-\epsilon$ elsewhere
1	1 at position of Player Two, $-\epsilon$ elsewhere
2	ply scale \cdot current ply $\cdot \epsilon$
3	ϵ if Player One is to move, else $-\epsilon$
4	ϵ (to indicate extent of board to convolution)

Figure 4: Board representation for the opposition game. The placeholder value $\epsilon \equiv \frac{1}{wh}$ is used to normalize the input.

assuming all possible outcomes are equally spaced (there are a finite number of these, as we impose a maximum number of ply on the game). All ply values are scaled by a factor of 0.1 in both input and output of the network to keep them in a reasonable range.

The board representation used for input to the network is specified in Figure 4.

The networks are trained on the 5 most recent generations of self-play games for 5 epochs with a constant learning rate of 0.005, using SGD with Nesterov momentum of 0.9. The loss is a linear combination of cross-entropy loss on the policy and result outputs and MSE loss on the value and plies-remaining outputs. For value heads, these coefficients are 20 for the policy outputs and 1 for the value output; for outcome heads, they are 100 for the policy output, 3 for the three result outputs, and 1 for the ply outputs. These coefficients were chosen to roughly balance the losses from the different categories of output nodes and tuned to work well in practice.

During self-play, AlphaZero-style Monte Carlo Tree Search is performed with $20h$ visits, using a temperature of 1.0 and Dirichlet noise of 0.5. At test time, both temperature and Dirichlet noise are set to 0.

References

Laterre, A.; Fu, Y.; Jabri, M. K.; Cohen, A.-S.; Kas, D.; Hajjar, K.; Dahl, T. S.; Kerkeni, A.; and Beguir, K. 2018. Ranked reward: Enabling self-play reinforcement learning for combinatorial optimization. *arXiv preprint arXiv:1807.01672*.

Morandin, F.; Amato, G.; Fantozzi, M.; Gini, R.; Metta, C.; and Parton, M. 2019. SAI: a sensible artificial intelligence that plays with handicap and targets high scores in 9x9 go (extended version). *CoRR* abs/1905.10863.

Ng, A. Y.; Harada, D.; and Russell, S. 1999. Policy invariance under reward transformations: Theory and application to reward shaping. In *ICML*, volume 99, 278–287.

Rechenberg, I. 1973. Evolution strategy: Optimization of technical systems by means of biological evolution. *Fromman-Holzboog, Stuttgart* 104:15–16.

Silver, D.; Hubert, T.; Schrittwieser, J.; Antonoglou, I.; Lai, M.; Guez, A.; Lanctot, M.; Sifre, L.; Kumaran, D.; Graepel, T.; Lillicrap, T. P.; Simonyan, K.; and Hassabis, D. 2017. Mastering chess and shogi by self-play with a general reinforcement learning algorithm. *CoRR* abs/1712.01815.

van Hasselt, H. P.; Guez, A.; Hessel, M.; Mnih, V.; and Silver, D. 2016. Learning values across many orders of magnitude. In *Advances in Neural Information Processing Systems*, 4287–4295.

Wierstra, D.; Schaul, T.; Glasmachers, T.; Sun, Y.; Peters, J.; and Schmidhuber, J. 2014. Natural evolution strategies. *The Journal of Machine Learning Research* 15(1):949–980.

Wirth, C.; Akrou, R.; Neumann, G.; and Fürnkranz, J. 2017. A survey of preference-based reinforcement learning methods. *The Journal of Machine Learning Research* 18(1):4945–4990.

Wu, D. J. 2019. Accelerating self-play learning in go. *CoRR* abs/1902.10565.

Zap, A.; Joppen, T.; and Fürnkranz, J. 2019. Deep ordinal reinforcement learning. *arXiv preprint arXiv:1905.02005*.